

PCD - H1 - Technical Report

Pasat Tudor Cosmin

1. Introduction

1.1. Homework objectives

1.2. Implementation requirements

1.3. Output requirements

2. Experiment

2.1. TCP

2.1.1. Streaming

2.1.2. Stop and wait

2.2. UDP

2.2.1. Streaming

2.2.2. Stop and wait

3. Discussion

3.1. Local host

3.2. EC2 Instance

4. Conclusion

5. References

1. Introduction

This technical report explores the performance analysis of data transfer protocols, with a focus on UDP (User Datagram Protocol) and TCP (Transmission Control Protocol). The primary objective is to measure the time required for transferring varying amounts of data across different message sizes and transmission mechanisms.

1.1. Homework objectives

The task entails designing a program capable of measuring data transfer time for different message sizes (1 to 65535 bytes) and amounts of data using UDP and TCP protocols. Additionally, the program must support both streaming and stop-and-wait transmission mechanisms.

1.2. Implementation requirements

- **Supported Protocols:** UDP and TCP protocols must be supported as parameters for both client and server components.
- **Message Size:** The program should facilitate data transfer across a range of message sizes to assess performance implications.
- **Transmission Mechanisms:** Two transmission mechanisms, streaming and stop-and-wait, should be implemented for comprehensive analysis.

1.3. Output requirements

- **Server Output:** The server will print the protocol used, number of messages read, and bytes read after each session.
- **Client Output:** Upon completion, the client will display transmission time, number of sent messages, and total bytes sent.

2. Experiment

The experiments were done on a local host, using a laptop with Intel Core I5 processor and MacOS Ventura 13.4.1. I ran all the pairs, protocol and mechanism, 100 times on local host and 10 times on an EC2 instance from AWS, and computed the minimum, maximum and the average values for each of the statistics presented in section 1.3.

For the stop and wait mechanism, I added to the client a timeout of 10 seconds. If he doesn't receive the acknowledgement until the timer runs out, the client will not move to the next package and will try to send the same one until the acknowledgement is received.

2.1. TCP

I used for both of the mechanisms a buffer of 65535 bytes.

2.1.1. Streaming

Type	Size	Messages			Bytes			Transmission time		
		Min	Avg	Max	Min	Avg	Max	Min	Avg	Max
LOCAL	500 MB	8001	8001.45	8036	524288000	524288000	524288000	0.18s	0.22s	0.5s
	1GB	16385	16385.13	16388	1073741824	1073741824	1073741824	0.36s	0.44s	0.97s
	2GB	32769	32769.39	32776	2147483648	2147483648	2147483648	0.63s	0.96s	2.64s
EC2	500 MB	24011	31204.7	44571	524288000	524288000	524288000	11.32s	13.60s	26.27s
	1GB	56770	86617.7	103460	1073741824	1073741824.0	1073741824	23.63s	26.53s	29.92s
	2GB	104484	153430.4	190772	2147483648	2147483648	2147483648	48.94s	54.13s	71.34s

Table 1: TCPS Client statistics

Type	Size	Messages			Bytes		
		Min	Avg	Max	Min	Avg	Max
LOCAL	500MB	8156	8289.31	9675	524288000	524288000	524288000
	1GB	16533	16727.15	17229	1073741824	1073741824	1073741824
	2GB	33121	33692.17	39712	2147483648	2147483648	2147483648
EC2	500MB	48620	67516.4	105008	524288000	524288000	524288000
	1GB	120228	193105.5	239137	1073741824	1073741824	1073741824
	2GB	222792	349565.4	435726	2147483648	2147483648	2147483648

Table 2: TCPS Server statistics

2.1.2. Stop and wait

Type	Size	Messages			Bytes			Transmission time		
		Min	Avg	Max	Min	Avg	Max	Min	Avg	Max
L O C A L	500 MB	8001	8001	8001	52428 8000	52428 8000	52428 8000	0.18s	0.20s	0.42s
	1GB	16385	16385	16385	10737 41824	10737 41824	10737 41824	0.38s	0.48s	1.29s
	2GB	32769	32769	32769	21474 83648	21474 83648	21474 83648	0.73s	1.08s	3.21s
E C 2	500 MB	8001	8001	8001	52428 8000	52428 8000	52428 8000	485.1 9s	487.5 4s	488.8 9s
	1GB	16385	16385	16385	10737 41824	10737 41824	10737 41824	1070. 96s	1093. 12s	1105. 93s

Table 3: TCPSW Client statistics

Type	Size	Messages			Bytes		
		Min	Avg	Max	Min	Avg	Max
L O C A L	500 MB	8070	8119.47	8240	524288000	524288000	524288000
	1GB	16542	16716.57	17210	1073741824	1073741824	1073741824
	2GB	33013	33374.60	34246	2147483648	2147483648	2147483648
E C 2	500 MB	58811	81002.33	116296	524288000	524288000	524288000
	1GB	105339	177299.66	219937	1073741824	1073741824	1073741824

Table 4: TCPSW Server statistics

2.2. UDP

I used for both of the mechanisms a buffer of 9216 bytes on the local host. For the tests on EC2, I tried using different values for the buffer, ranging from 1000 to 65535 and finally I found that 1000 is the value that is the best to use in order to minimise the number of data loss, on the streaming mechanism and 60000 on stop and wait.

2.2.1. Streaming

Type	Size	Messages			Bytes			Transmission time		
		Min	Avg	Max	Min	Avg	Max	Min	Avg	Max
L O C A L	500 MB	56889	56889	56889	524288000	524288000	524288000	0.76	1.16	0.92
	1GB	116509	116509	116509	1073741824	1073741824	1073741824	1.66s	1.91s	3.45s
	2GB	233017	233017	233017	2147483648	2147483648	2147483648	3.09s	3.76s	7.89s
E C 2	500 MB	524288	524288	524288	524288000	524288000	524288000	14.85s	15.07s	15.60s
	1GB	1073742	1073742	1073742	1073741824	1073741824	1073741824	30.49s	30.64s	30.86s
	2GB	2147484	2147484	2147484	2147483648	2147483648	2147483648	61.14s	63.09s	64.61s

Table 5: UDPS Client statistics

Type	Size	Messages			Bytes		
		Min	Avg	Max	Min	Avg	Max
L O C A L	500M B	56738	56883.16	56889	522896384	524234178.56	524288000
	1GB	116047	116501.6 4	116509	106948403 2	1073673994.2 4	107374182 4
	2GB	231733	232985.9 5	23301 7	213565030 4	2147197491.2	214748364 8
E C 2	500M B	509366	512014.4	51483 3	509366000	512014400.0	514833000
	1GB	865633	926466.2	98794 6	865632824	926466024.0	987945824
	2GB	189184 4	1935846	21096 22	189184364 8	1935845648	210962164 8

Table 6: UDPS Server statistics

2.2.2. Stop and wait

Type	Size	Messages			Bytes			Transmission time		
		Min	Avg	Max	Min	Avg	Max	Min	Avg	Max
L O C A L	500 MB	56889	56889	56889	52428 8000	52428 8000	52428 8000	2.02s	2.18s	3.11s
	1GB	116509	116509	116509	10737 41824	10737 41824	10737 41824	4.18s	4.53s	7.35s
	2GB	233017	233017	233017	21474 83648	21474 83648	21474 83648	8.35s	9.37s	13.34 s
E C 2	10 MB	274	280.66	290	16397 280	16806 773.33	17371 520	190.7 9s	202.1 2s	218s

Table 7: UDPSW Client statistics

Type	Size	Messages			Bytes		
		Min	Avg	Max	Min	Avg	Max
L O C A L	500MB	56889	56889	56889	524288000	524288000	524288000
	1GB	116509	116509	116509	1073741824	1073741824	1073741824
	2GB	233017	233017	233017	2147483648	2147483648	2147483648
E C 2	10MB	176	176	176	10531520	10531520	10531520

Table 8: UDPSW Server statistics

3. Discussion

3.1. Local host

We can see from the tables above that the UDP protocol, using the streaming mechanism, is the worst when considering the number of bytes that actually get to the server. There are packages of bytes lost even for the smaller amounts of data.

As for the other 3, all the data gets to the server, but the transmission times are significantly different, between the 2 TCP programs and the UDP one. This is mostly because I used a much smaller buffer size for the UDP, in order to avoid getting more expensive getting lost.

The TCP with stop and wait mechanism has a higher latency than the streaming version, because the client needs to wait for acknowledgement before sending the next package.

3.2. EC2 Instance

- TCP with Streaming - unlike the other methods, just worked when I run it on cloud, without any modification. We can see from **Table 1**, that compared to the local host, the transmission time is significantly higher. This happens because the communication goes through the internet and there are several

more factors that intervene, like the reliability of the network, the bandwidth, which is never unlimited, the latency, which is never 0, etc.

- TCP with Stop and Wait - for this mechanism I modified the code a little bit, because socket library implements the TCP communication and write to a buffer from where we, as users of this library, read the messages, we can get just a piece of the package in one message, the rest being read the next iteration. Because of this, I added to send the acknowledgement only when the entire package is sent, which, we can see from **Table 3**, affects the transmission time, going from 13 seconds on average for the streaming mechanism, to 487 seconds, for this one, on 500MB of data. This shows that the latency is a lot higher for this one, proving again that it is never 0.
- UDP with Streaming - because the bandwidth isn't unlimited for the network communication with the cloud, I had to add a sleep for the UDP with the streaming mechanism version of 0.1 seconds, every 6000 packages sent, otherwise an exception is thrown, because the buffer is full. Also, as I mentioned before, I had to use a package size of only 1000 bytes, in order to assure that the data loss was minimal. Although this configuration increases the transmission time, because the number of packages is higher and the time spent sleeping. In the end, the transmission time is actually a little higher than the one of TCP streaming and we don't have the reliability of TCP
- UDP with Stop and Wait - this mechanism brought the most problems. Since it's a UDP connection, there is a probability to lose both the packages sent by the client and the acknowledgment sent back by the server and at some point the client and the server are both stuck waiting for messages from each other. To solve this I added a timeout on the client side for it to receive the acknowledgement. If it expires it tries to send the same package again. Also on the server side I computed the hash of the current package received and compared it to the last one. If they are the same it assumes that the acknowledgement was lost and the same package was sent, and just sends the same message back, without saving the package. Since there is this retransmission method, I tried using packages of different sizes, finally using

the 60000 value. Unfortunately, all these new steps that need to be done increase the transmission very much, taking 200 seconds to send only 10MB of data. Most of the time is spent waiting for responses that never actually come and end up with a timeout exception and retransmission of the package.

4. Conclusion

When using a remote server, there are fallacies of distributed computing that need to be taken into consideration. As we saw from the experiments, we can't assume the bandwidth is unlimited or that the latency is zero. Also on the local host I was the only administrator, meanwhile when using the cloud, the data goes through different subnets, until it gets to my implemented server. Also, I can only control the network topology in my home, outside there can always be changes in the structure of the network and as a result change the experiment's outcomes.

In the end, we can say that one protocol or mechanism is better than another. All have their advantages and disadvantages, and we need to analyse the use case to find which is best suited for it.

5. References

- <https://profs.info.uaic.ro/~adria/teach/courses/pcd>
- https://profs.info.uaic.ro/~adria/teach/courses/pcd/homework/PCD_Homework1.pdf
- <https://docs.python.org/3.11/library/socket.html#module-socket>
- <https://pythontic.com/modules/socket/introduction>
- <https://pythontic.com/modules/socket/udp-client-server-example>
- <https://forums.developer.apple.com/forums/thread/74655#:~:text=By%20default%20macOS%20has%20limited,following%20command%20in%20the%20terminal.>